

Back to Basics: Smart Pointers

Arthur O'Dwyer
2019-09-19

Outline

- Unique ownership with `std::unique_ptr` [5–18]
- Shared ownership with `std::shared_ptr` [19–30]
- `make_shared` and `make_unique` [31–39]
- `std::weak_ptr` [40–46]
- `std::enable_shared_from_this` [47–52]
- Questions?

C++11 (and up) smart pointer types

auto_ptr

C++98. Deprecated in C++11. Removed in C++17.

unique_ptr

C++11 replacement for auto_ptr. C++14 adds make_unique.

shared_ptr

C++11. Reference-counting. C++17 adds shared_ptr<T[]>.

weak_ptr

C++11. "Weak" references.

All the standard smart pointers, plus make_shared<T>, arrived together in C++11.

make_shared<T[]> finally arrives in C++20.

Smart pointers look familiar, by design

```
std::shared_ptr<String> p =  
    std::make_shared<String>("Hello");  
  
auto q = p;  
  
p = nullptr;  
  
if (q != nullptr) {  
    std::cout << q->Length() << *q << '\n';  
}
```


Creation of
heap-allocated objects
is slightly novel
(more on this later)

Initialization (and
assignment) works as
you'd expect

So does `nullptr`

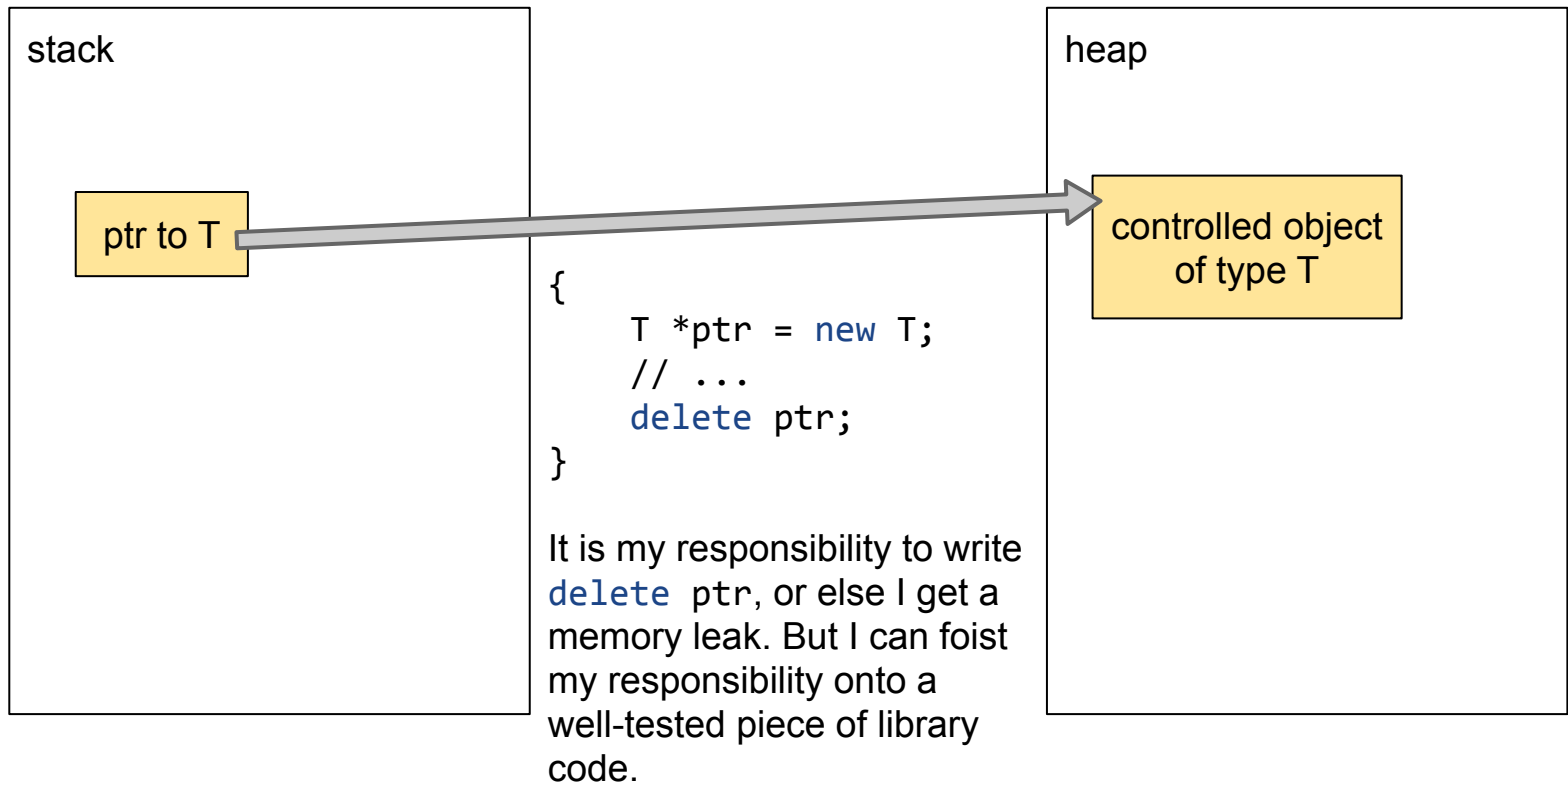
So does comparison
and dereferencing

Scott Meyers EMC++ Item 18

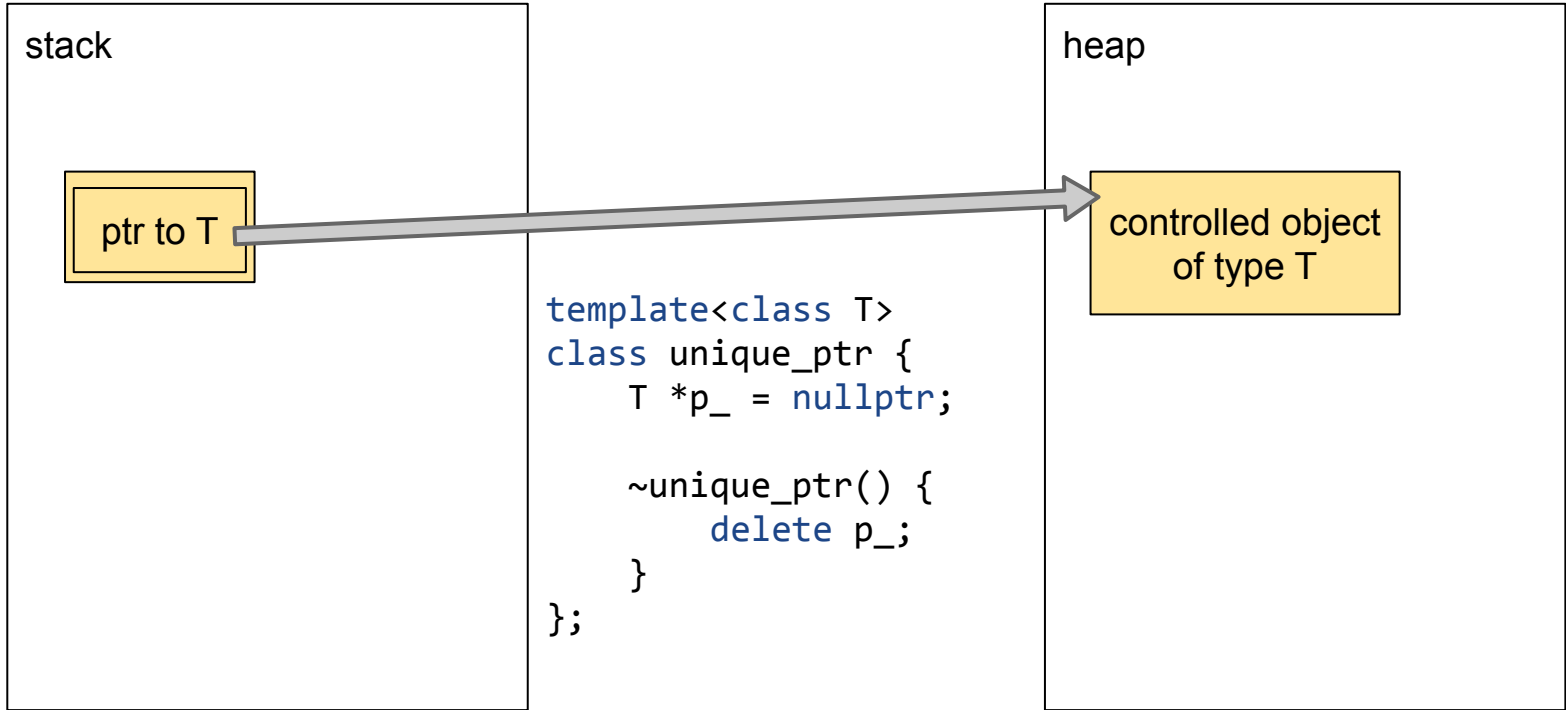


Use `std::unique_ptr` for
exclusive-ownership
resource management.

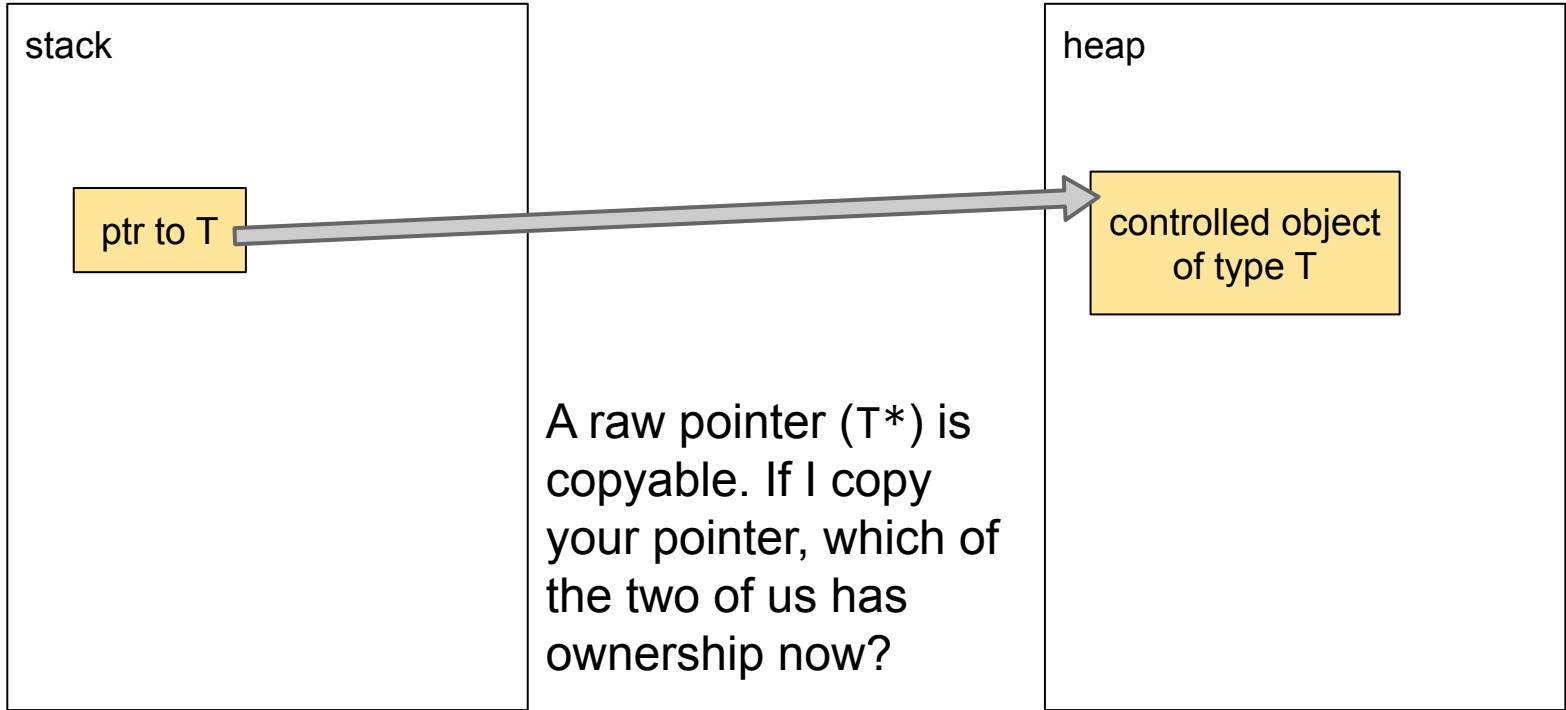
Exclusive ownership



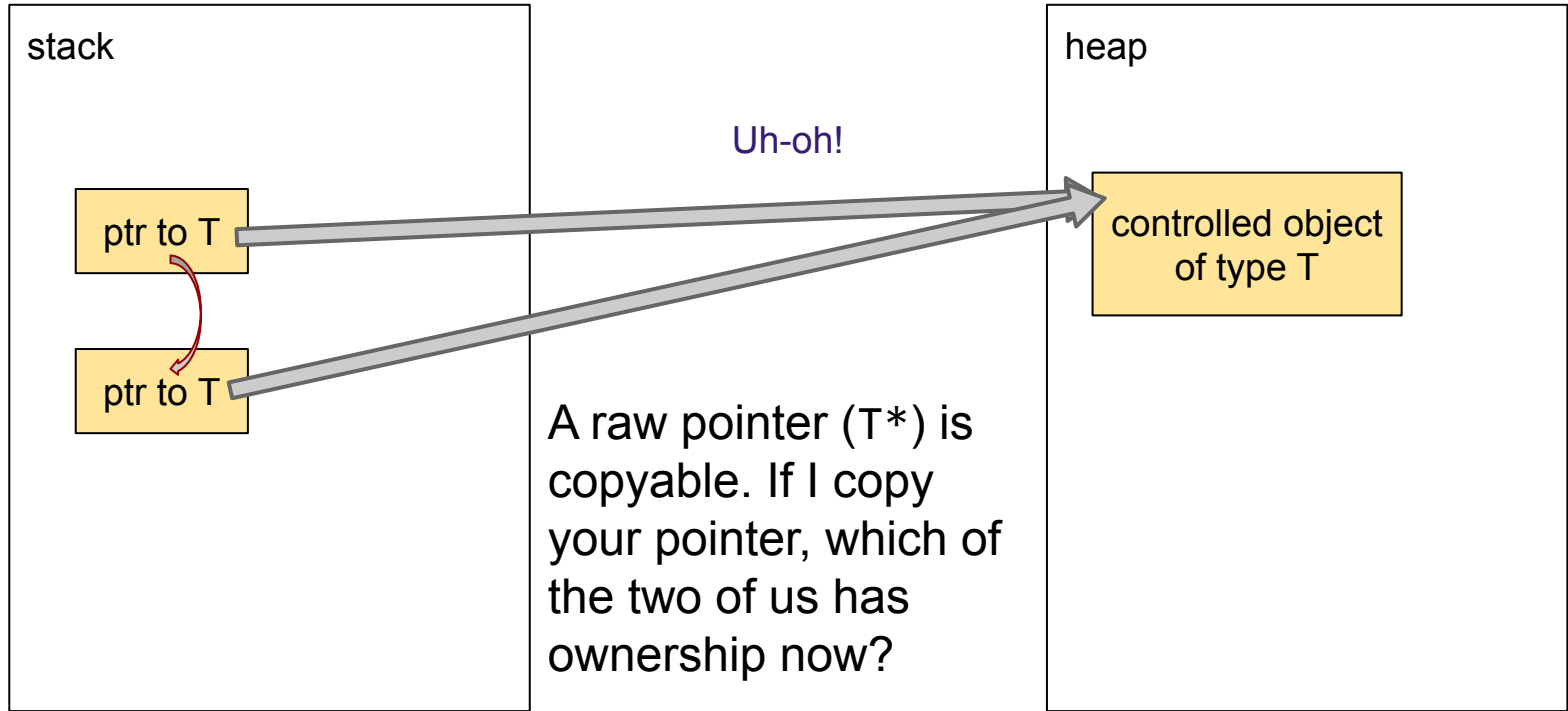
std::unique_ptr<T>




unique_ptr is moveable (move-only)




unique_ptr is moveable (move-only)



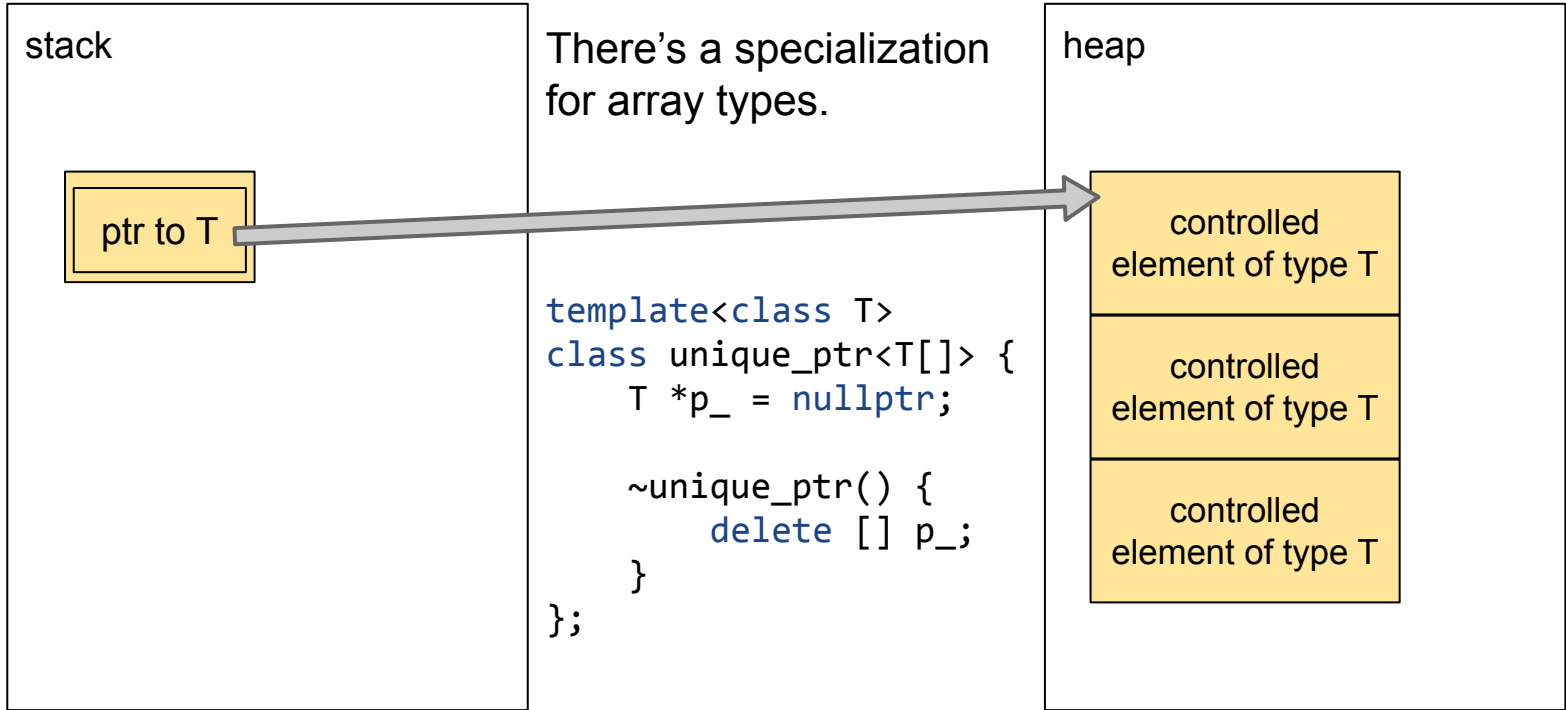
unique_ptr is moveable (move-only)



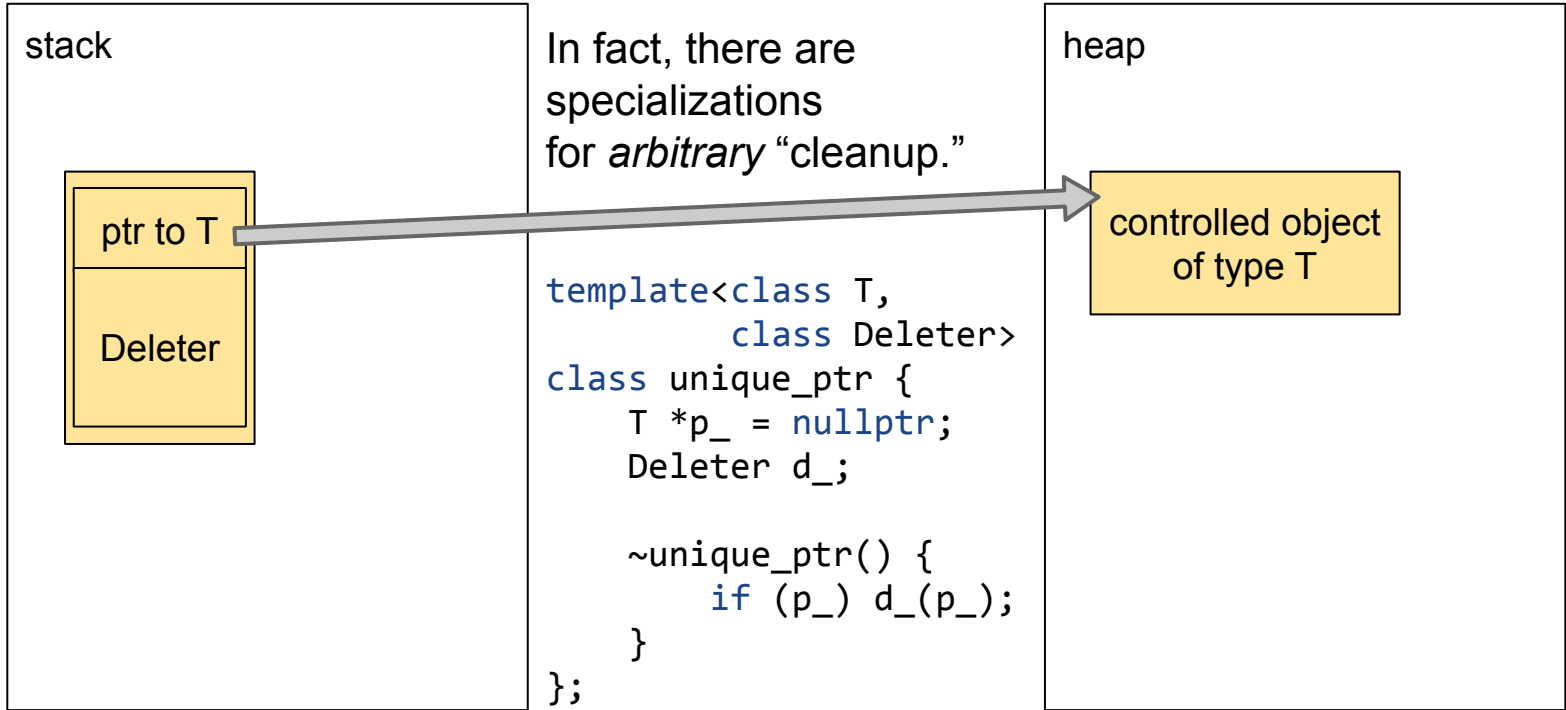
unique_ptr is moveable (move-only)



std::unique_ptr<T[]>



std::unique_ptr<T, Deleter>



std::unique_ptr<T, Deleter>

```
template<class T, class Deleter = std::default_delete<T>>
```

```
class unique_ptr {  
    T *p_ = nullptr;  
    Deleter d_;  
  
    ~unique_ptr() {  
        if (p_) d_(p_);  
    }  
};
```

```
template<class T>  
struct default_delete {  
    void operator()(T *p) const {  
        delete p;  
    }  
};
```

unique_ptr is always a template of two parameters. If you provide no second parameter, it is defaulted to std::default_delete<T>.

Even this part is customizable. If Deleter::pointer names a type, then this member will be of that type instead of T*.

Custom deleters can do neat things

```
struct FileCloser {  
    void operator()(FILE *fp) const {  
        assert(fp != nullptr);  
        fclose(fp);  
    }  
};
```

```
FILE *fp = fopen("input.txt", "r");  
std::unique_ptr<FILE, FileCloser> uptr(fp);
```

Custom deleters can do neat things

My codebase's OpenSSL code is peppered with this pattern:

```
struct REQDeleter {  
    void operator()(X509_REQ *p) const { X509_REQ_free(p); }  
};  
  
class MyCSR {  
    std::unique_ptr<X509_REQ, REQDeleter> p_;  
public:  
    explicit MyCSR(std::unique_ptr<X509_REQ, REQDeleter> p) :  
        p_(std::move(p)) {}  
};
```


Custom deleters can do neat things

```
explicit MyCSR(std::unique_ptr<X509_REQ, REQDeleter> p) :  
    p_(std::move(p)) {}
```

This constructor takes a `unique_ptr` by value.

Functions taking owning pointers by value are sometimes called ***sinks***.

A “`unique_ptr` sink” clearly conveys ***transfer of ownership***.

To call this constructor at all, you must have a `unique_ptr` already — so, you must have ownership of an `X509_REQ` structure. By taking your `unique_ptr` by value (by move), the constructor shows that it ***takes*** ownership of that `X509_REQ` structure away from you.


`unique_ptr` is the glue between the low-level, non-RAII, raw resource (`X509_REQ`) and the high-level RAII business object (`MyCSR`).

Rules of thumb for smart pointers

- Treat smart-pointer types just like raw pointer types.
 - Pass by value.
 - Return by value (of course).
 - Passing a pointer by reference is usually a code smell.
 - Same goes for smart pointers.
- A function taking a `unique_ptr` by value shows ***transfer of ownership***.
 - Even shows exactly ***what*** responsibility is being transferred, because the responsibility is encoded in the deleter type.
 - Usually the responsibility is simply “to call `delete`.”
- Smart pointers are frequently implementation details and glue.
 - To bake `unique_ptr` or `shared_ptr` into your ***interface*** might be a code smell. Try to deal in business classes like `MyCSR` instead.

EMC++ Item 19

Use `std::shared_ptr` for
shared-ownership
resource management.



`std::shared_ptr<T>`

`shared_ptr` looks similar to `unique_ptr` on the surface...

```
std::unique_ptr<int> uptr = std::make_unique<int>(42);  
std::shared_ptr<int> sptr = std::make_shared<int>(42);
```

But it is vastly more complicated on the inside!

`shared_ptr` expresses shared ownership. Specifically, ***reference-counting***.

Simple reference counting

“Will the last person out of the room please turn out the lights.”

How do you know if you're the last person in the room? (It's a very big room.)

We keep a jar of tokens by the door. When anyone enters, they put a token in the jar. When anyone leaves, they take a token with them. If you take the ***last*** token from the jar, then you must be the last person to leave the room.

Notice that everyone involved must cooperate in our scheme. Anyone who enters without leaving a token in the jar, risks being left in the dark!

Simple reference counting

```
struct Widget {  
    std::atomic<int> usecount_{1};  
    Widget *retain() { ++usecount_; return this; }  
    void release() { if (--usecount_ == 0) delete this; }  
};  
  
{  
    Widget *p = new Widget;  
    Widget *q = p->retain();  
    p->release();  
    q->release(); // causes the Widget to be deleted  
}
```

shared_ptr automates refcounting

```
struct Widget {  
    std::atomic<int> usecount_{1};  
    Widget *retain() { ++usecount_; return this; }  
    void release() { if (--usecount_ == 0) delete this; }  
};  
  
{  
    Widget *p = new Widget;  
    Widget *q = p->retain();  
    p->release();  
    q->release();  
}
```

`shared_ptr<Widget>` will have a copy-constructor that increments the refcount, and a destructor that decrements the refcount.

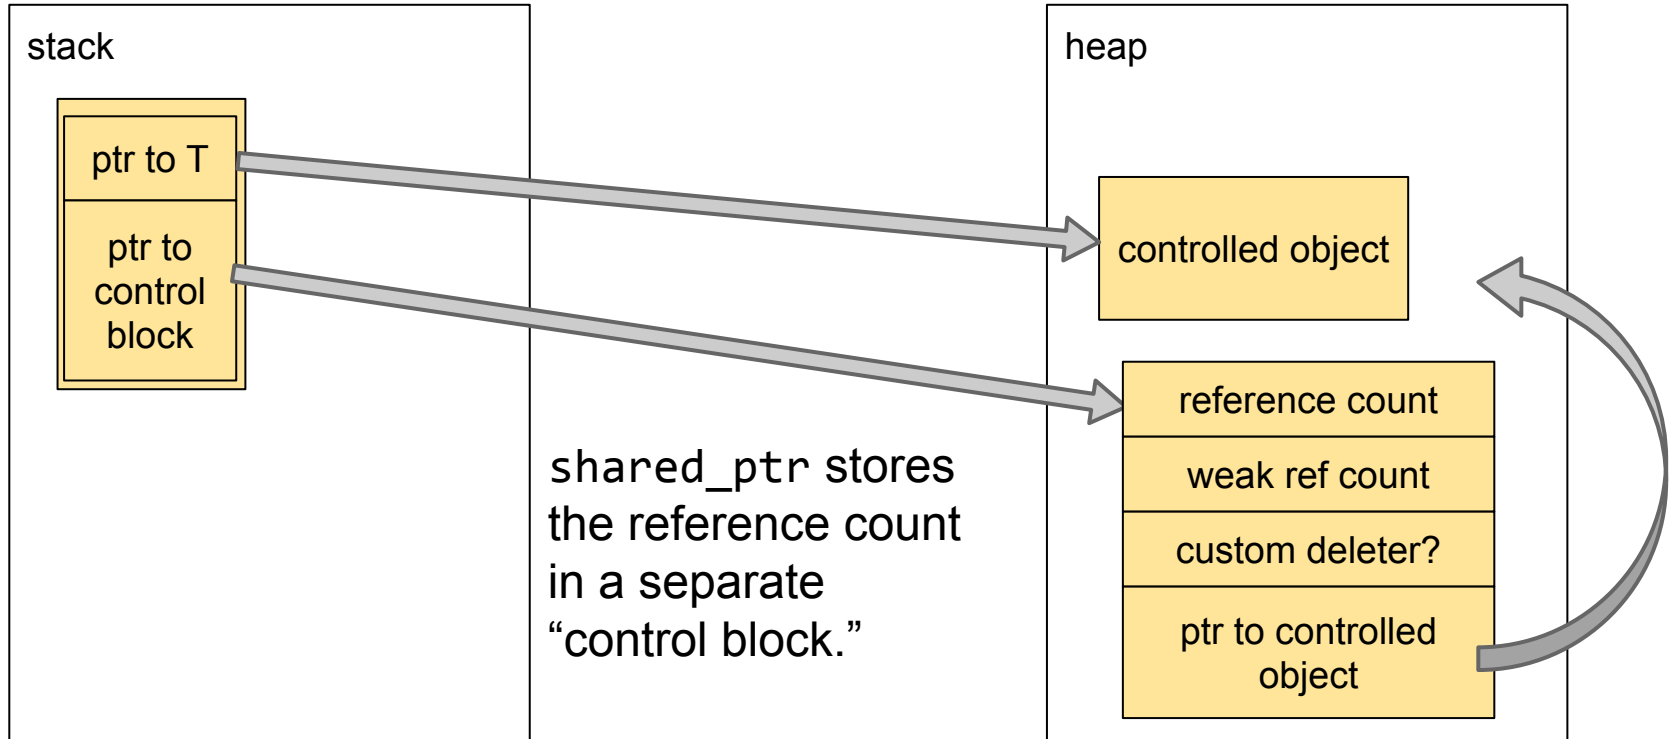
But there's a problem...

C++ objects don't contain refcounts!

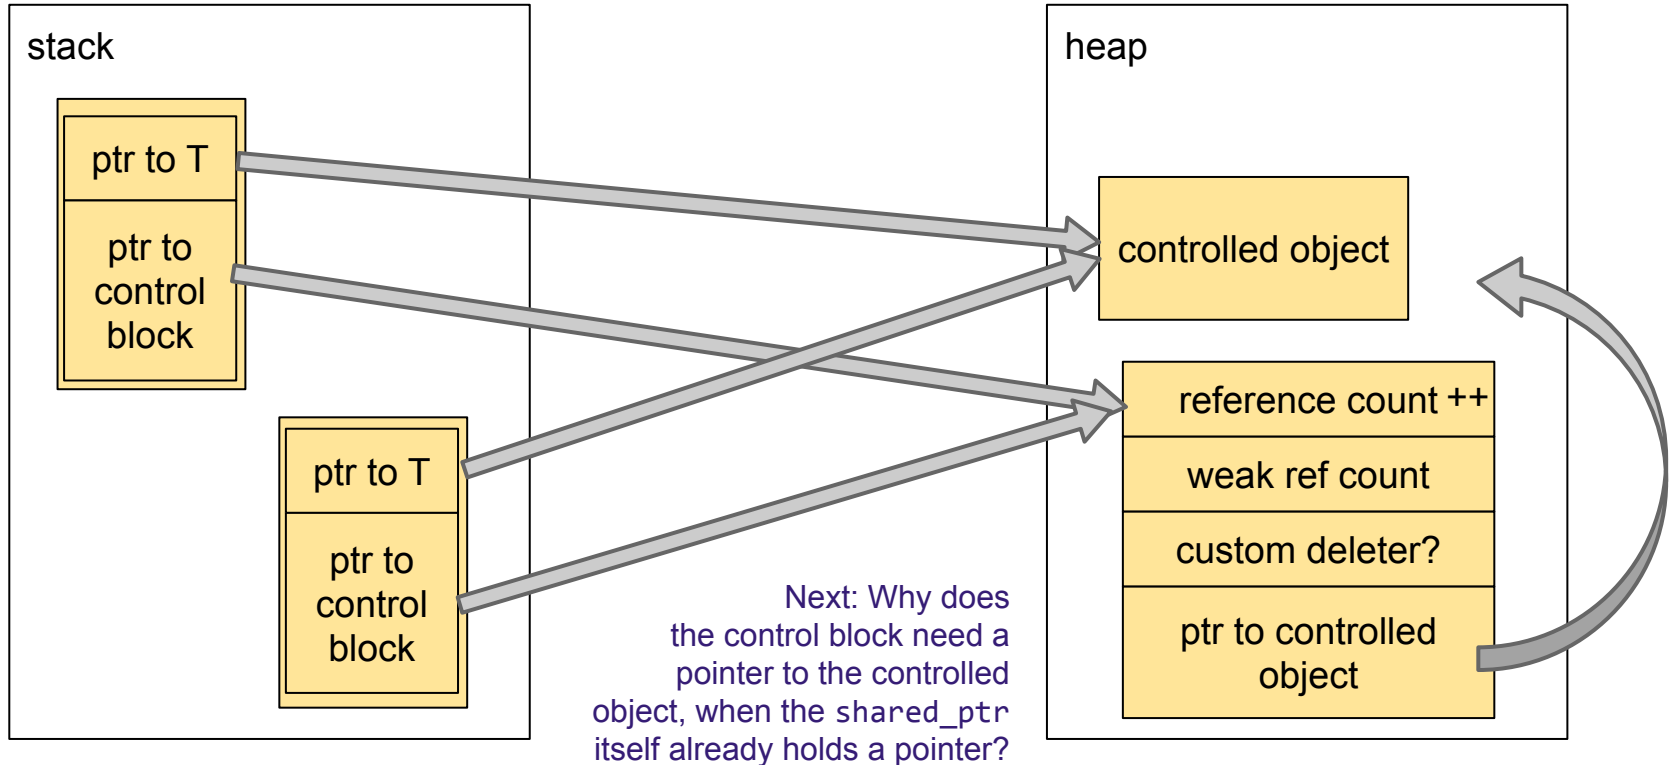
```
struct Widget {  
    std::atomic<int> usecount_{1};  
    Widget *retain() { ++usecount_; return this; }  
    void release() { if (--usecount_ == 0) delete this; }  
};  
  
{  
    Widget *p = new Widget;  
    Widget *q = p->retain();  
    p->release();  
    q->release();  
}
```

What about
`shared_ptr<int>`? Where
could we store the reference
count for an `int` object?

`std::shared_ptr<T>`



Copying a `std::shared_ptr`





Detour: Physical class layout


```
struct Fruit { int juice; };
```

```
struct Vegetable { int fiber; };
```

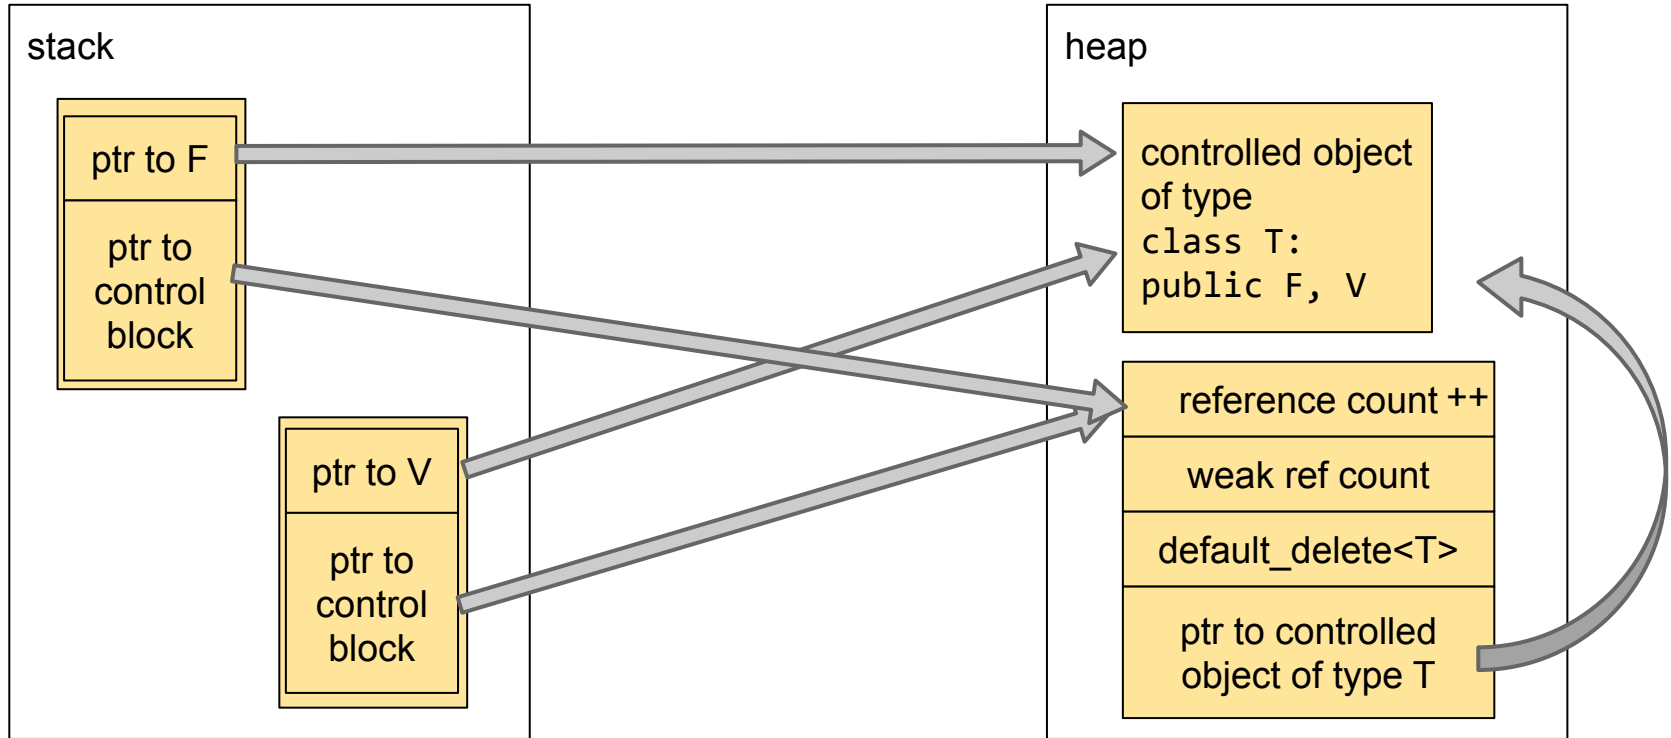
```
struct Apple : Fruit { int redness; };
```



```
struct Tomato : Fruit, Vegetable { int sauce; };
```



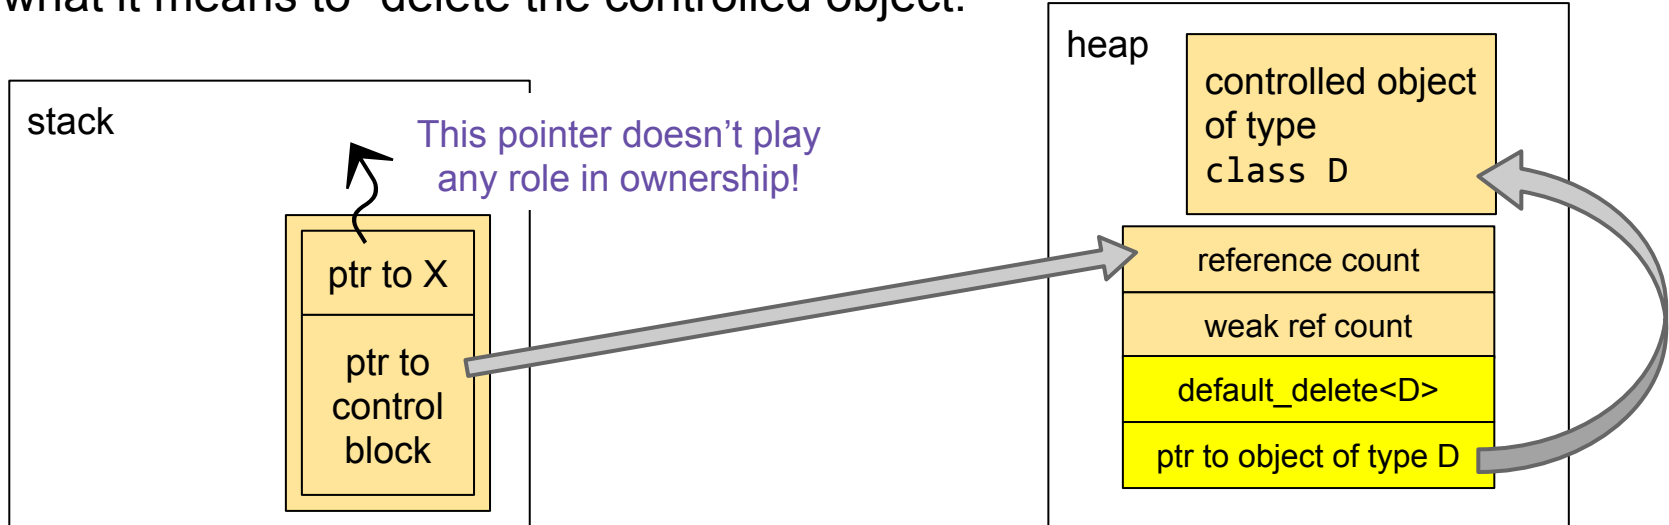
std::shared_ptr to base class



shared_ptr type-erases the deleter

shared_ptr, unlike unique_ptr, places a layer of indirection between the physical heap-allocated object and the notion of ownership.

Your shared_ptr instances are essentially participating in ref-counted ownership of the **control block**. The control block itself is sole arbiter of what it means to “delete the controlled object.”



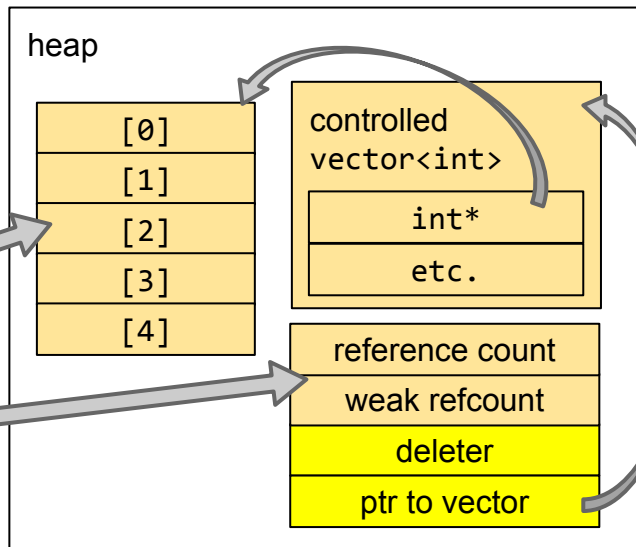

shared_ptr's “aliasing constructor”

```
using Vec = std::vector<int>;
```


```
std::shared_ptr<int> foo() {  
    auto elts = { 0,1,2,3,4 };  
    std::shared_ptr<Vec> pvec = std::make_shared<Vec>(elts);  
    return std::shared_ptr<int>(pvec, &(*pvec)[2]);  
}
```

Share ownership with pvec
but point to &(*pvec)[2]

```
int main() {  
    std::shared_ptr<int> ptr = foo();  
    for (int i = -2; i < 3; ++i) {  
        printf("%d\n", ptr.get()[i]);  
    }  
}
```



EMC++ Item 21



Prefer `std::make_unique`
and `std::make_shared`
to direct use of `new`.

Goal: No raw new or delete

The first rule of C++ memory management is: Every new must have a matching delete, and vice versa.

The second rule of C++ memory management is: Stop using manual calls to delete! Use smart pointers whose destructors automatically call delete at the right time.

But if we stop writing delete... by the first rule, mustn't we stop writing new?

```
auto *w = new Widget();  
use(w); // Is this code OK?
```

```
std::shared_ptr<Widget> w(new Widget());  
use(w); // Is this code any better?
```


Goal: No raw new or delete

Yes! If we're not going to write `delete`, we *should* stop writing `new`.

Whenever we heap-allocate something, we'll use a factory function that returns it *already wrapped in a smart pointer*.


We should never have to touch raw pointers with our hands.

```
auto w = std::make_shared<Widget>();  
use(w); // Now we have no reason  
        // to flag this code.  
        // No new, no delete.
```




make_shared can even be optimized

```
template<class T, class... Args>
std::shared_ptr<T> make_shared(Args&&... args) {
    T *rawptr = new T(std::forward<Args>(args)...);
    return std::shared_ptr<T>(rawptr); // take ownership from rawptr
}
```




make_shared can even be optimized

```
template<class T, class... Args>
std::shared_ptr<T> make_shared(Args&&... args) {
    T *rawptr = new T(std::forward<Args>(args)...);
    return std::shared_ptr<T>(rawptr); // take ownership from rawptr
}
```




make_shared can even be optimized

```
template<class T, class... Args>
std::shared_ptr<T> make_shared(Args&&... args) {
    T *rawptr = new T(std::forward<Args>(args)...);
    return std::shared_ptr<T>(rawptr); // take ownership from rawptr
}
```



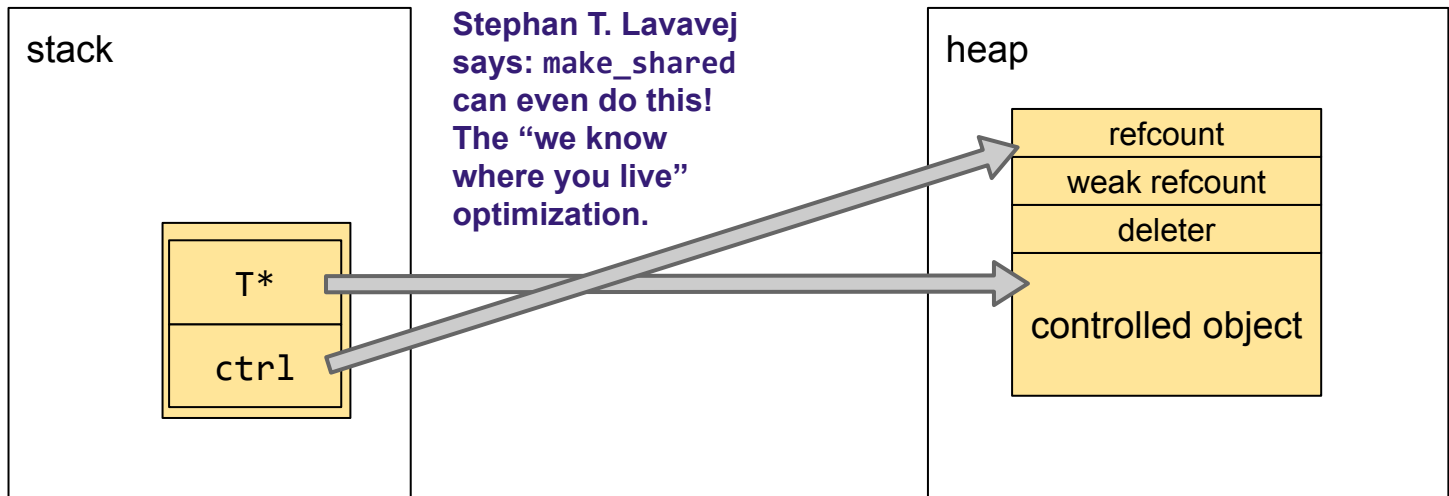
make_shared can even be optimized

```
template<class T, class... Args>
std::shared_ptr<T> make_shared(Args&&... args) {
    auto *rawptr = new ControlBlockAnd<T>(std::forward<Args>(args)...);
    return std::shared_ptr<T>::From(rawptr);
}
```



make_shared can even be optimized

```
template<class T, class... Args>
std::shared_ptr<T> make_shared(Args&&... args) {
    auto *rawptr = new ControlBlockAnd<T>(std::forward<Args>(args)...);
    return std::shared_ptr<T>::From(rawptr);
}
```



Conclusion: Use `make_shared`


- `make_shared` and `make_unique` wrap `raw new`, just as `~shared_ptr` and `~unique_ptr` wrap `raw delete`.
- If you ***never*** touch raw pointers with your hands, then you never need to worry about leaking them.
- `make_shared` can be a performance optimization!

By the way, `unique_ptr<T>` is implicitly convertible to `shared_ptr<T>...`

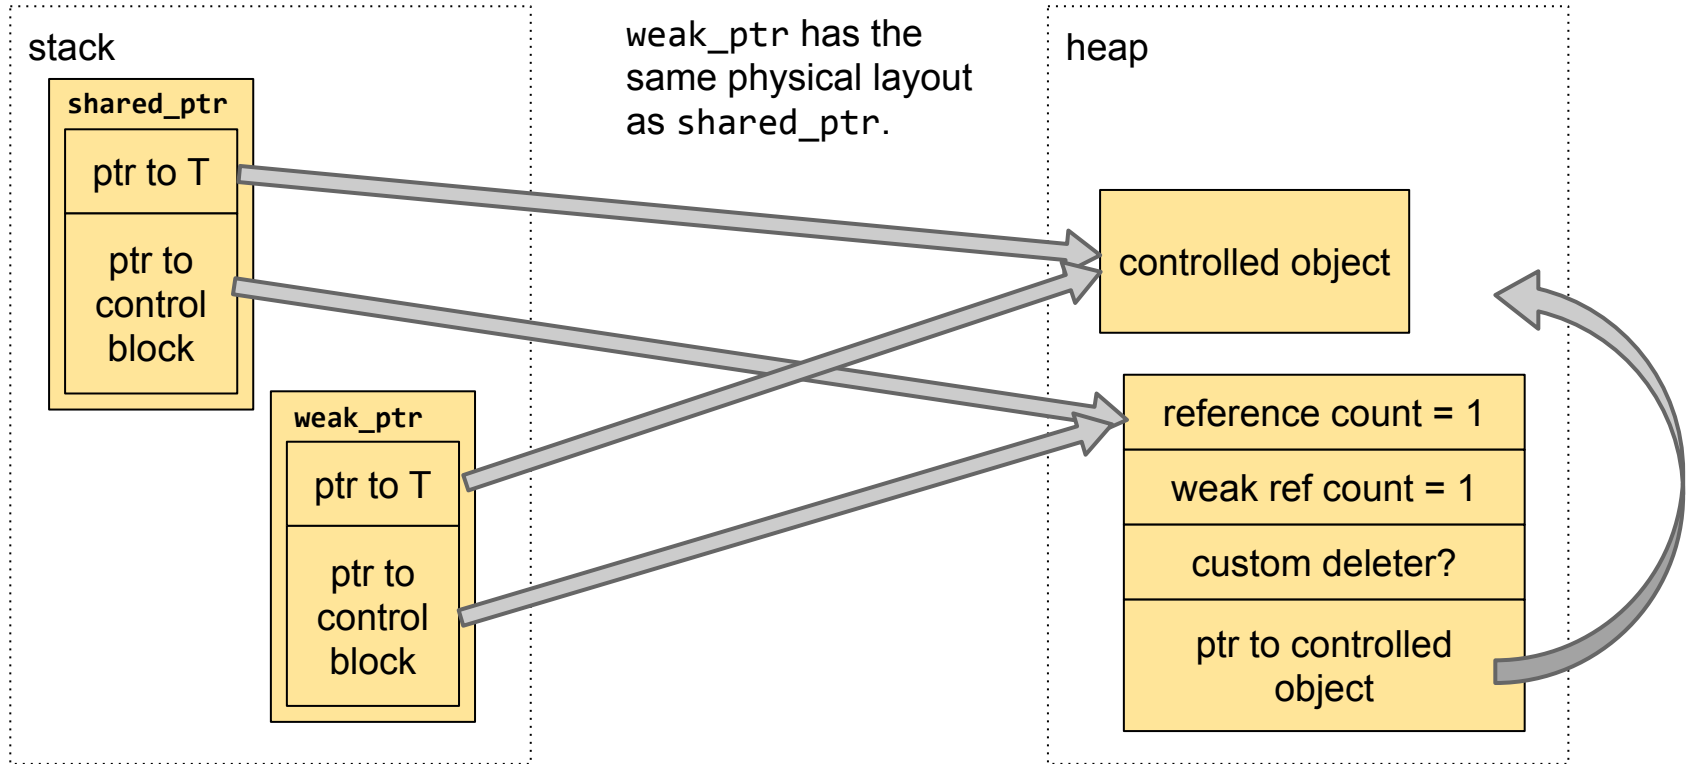
```
std::shared_ptr<Widget> sptr = std::make_unique<Widget>();
```

EMC++ Item 20

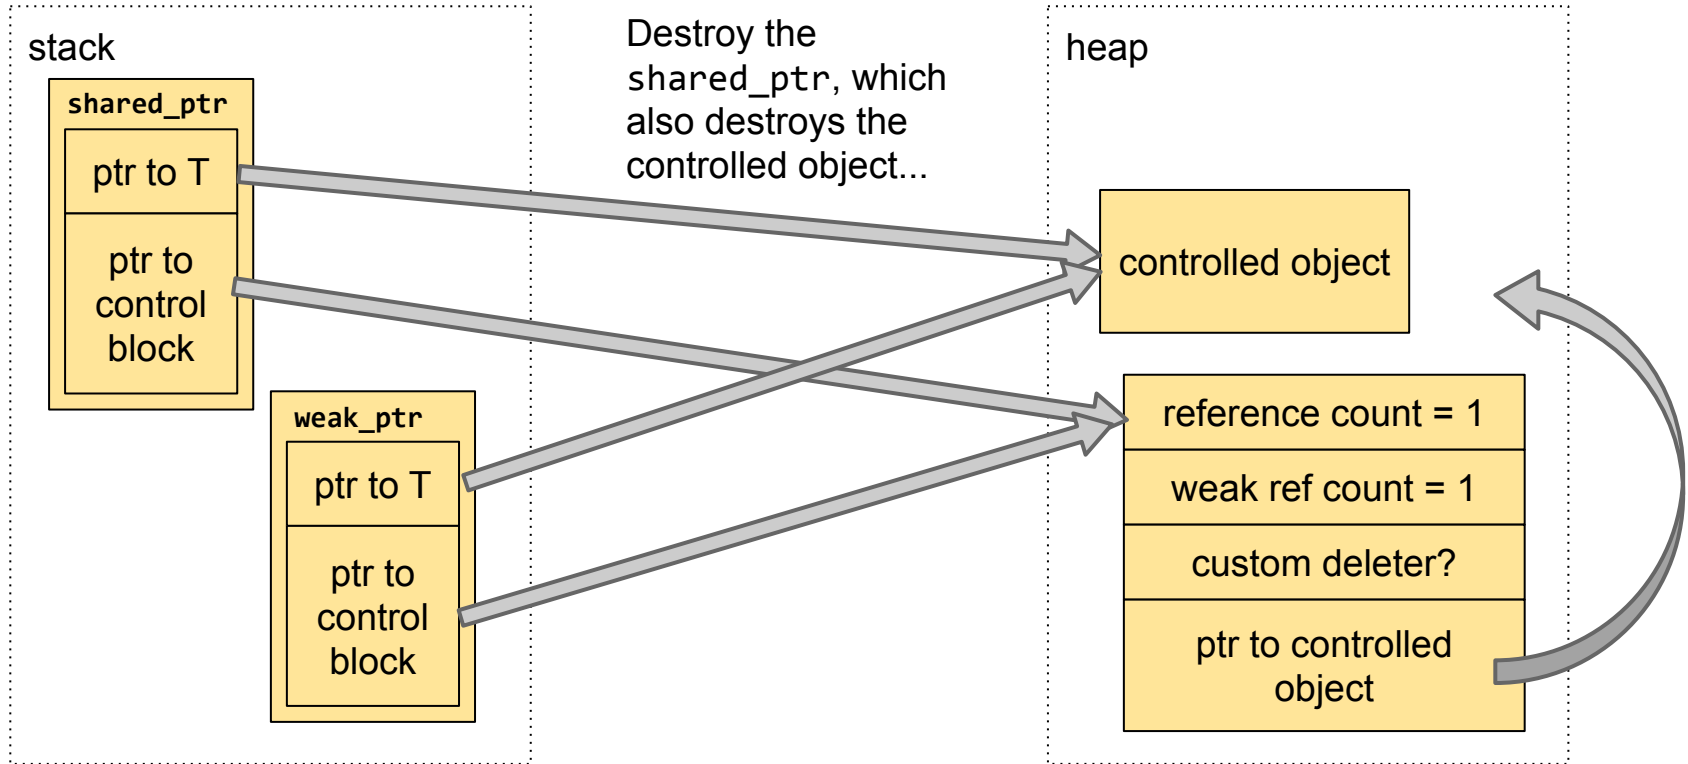
Use `std::weak_ptr` for
`shared_ptr`-like pointers
that can dangle
(*and know when they are dangling*)



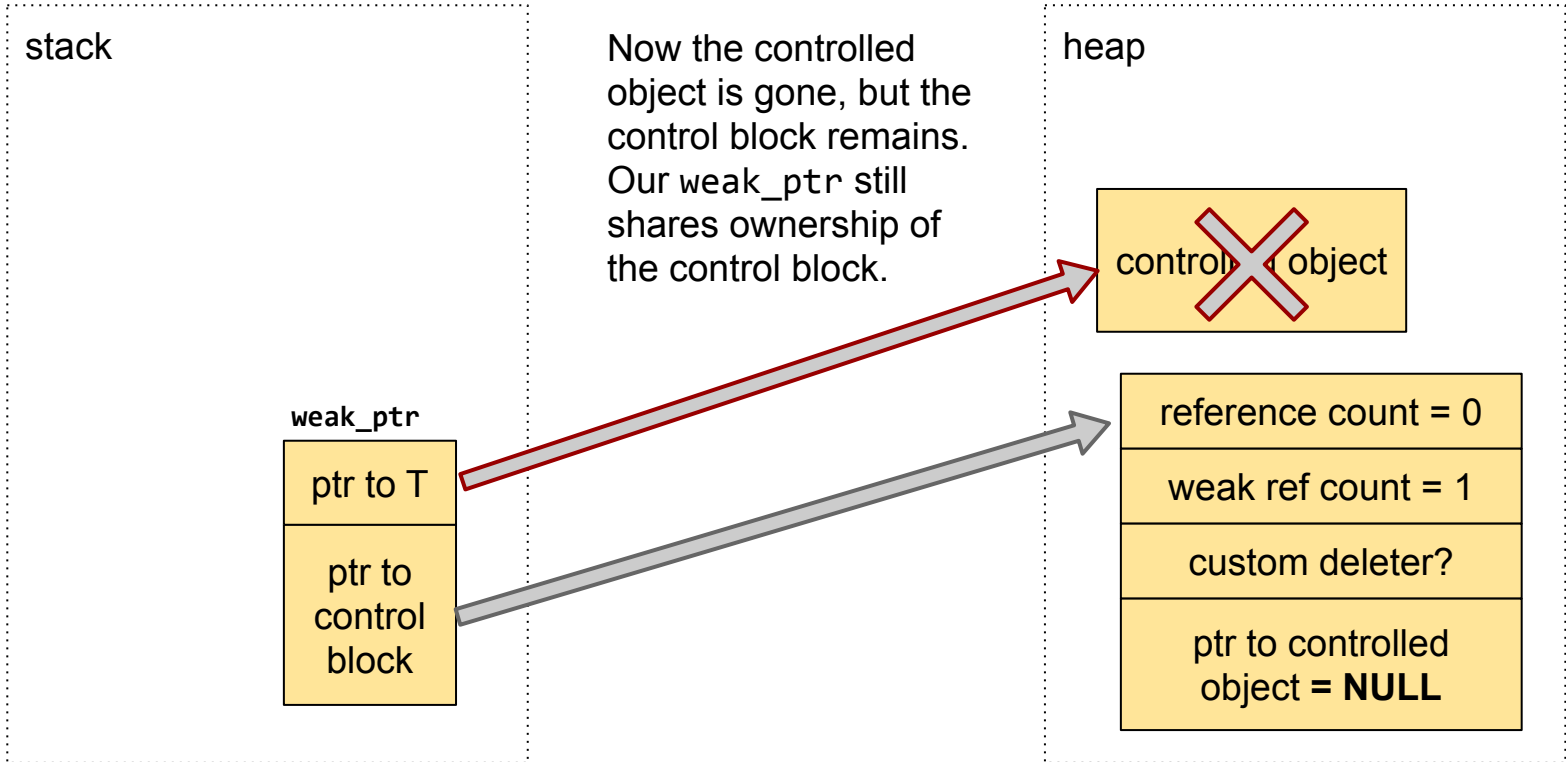
std::weak_ptr<T>



std::weak_ptr<T>



std::weak_ptr<T>



You can't dereference a weak_ptr

You can only convert it to a shared_ptr.

- Think of a weak_ptr as a “ticket for a shared_ptr.”
 - If you hold a weak_ptr, you're entitled to receive a shared_ptr (if the controlled object still exists at all, of course).
 - If you hold a shared_ptr, you're certainly entitled to get a weak_ptr!
- The “redeem a ticket” operation can be spelled in two ways: explicit type-conversion, or wptr.lock().
- The “get a ticket” operation can be spelled only as an explicit conversion. (Consider writing your own unlock() function?)

You can't dereference a weak_ptr

You can only convert it to a shared_ptr.

```
void recommended(std::weak_ptr<T> wptr) {
    std::shared_ptr<T> sptr = wptr.lock();
    if (sptr != nullptr) {
        use(sptr);
    }
}

void not_recommended(std::weak_ptr<T> wptr) {
    try {
        std::shared_ptr<T> sptr { wptr }; // call the explicit constructor
        use(sptr);
    } catch (const std::bad_weak_ptr&) {}
}
```

Idiomatic “redeem a ticket” one-liner

C++ has always supported variable declarations inside `if` conditions. This is one of the rare times it's a good idea to use that feature.

```
if (auto sptr = wptr.lock()) {  
    use(sptr);  
}
```

The other case is:

```
if (RedWidget *p = dynamic_cast<RedWidget*>(widgetptr)) {  
    use_red_widget(*p);  
}
```

What if a Widget is its own ticket?

Earlier, I said...

- Think of a `weak_ptr` as a “ticket for a `shared_ptr`.”
 - If you hold a `weak_ptr`, you’re entitled to receive a `shared_ptr` (if the controlled object still exists at all, of course).
 - If you hold a `shared_ptr`, you’re certainly entitled to get a `weak_ptr`!


What if you want the rule to be “If you hold a ***raw*** pointer, you’re entitled to receive a `shared_ptr`”?

Can we make that work?

What if a Widget is its own ticket?

Can we make that work? Yes. We just need to store a ticket (weak_ptr) somewhere that the Widget itself can get to it.

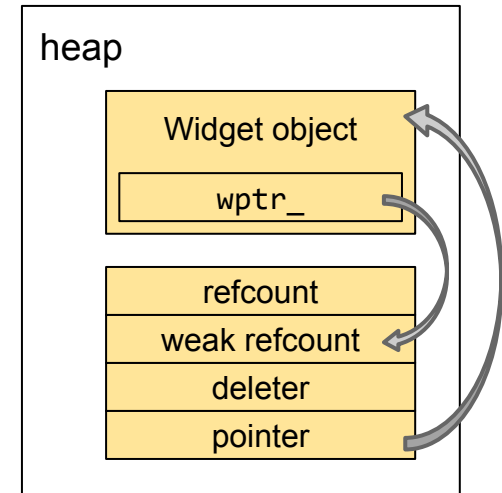
```
class Widget {  
    std::weak_ptr<Widget> wptr_ = ???;  
public:  
    std::shared_ptr<Widget> shared_from_this() {  
        return wptr_.lock();  
    }  
};
```



What if a Widget is its own ticket?

Pull out the weak_ptr into a special base class which is known to the STL implementation. Whichever constructor of shared_ptr first creates the control block also fills in the weak_ptr.

```
class Widget :  
    public std::enable_shared_from_this<Widget> {  
    ...  
};  
  
Widget *p = new Widget();  
assert(p->shared_from_this() == nullptr);  
  
std::shared_ptr<Widget> s1(p);  
assert(p->shared_from_this() == s1);
```



Curiously recurring template pattern

Notice that `enable_shared_from_this`:

- Cannot be implemented “in user-space,” because it must conspire with the implementation of `shared_ptr`’s constructor
- Uses the Curiously Recurring Template Pattern (CRTP)
 - The interface’s signature uses type `Widget`, so `Widget` must be a parameter to the base class template — even though `Widget` **derives from** that base class!

```
template<class T>
class enable_shared_from_this {
    shared_ptr<T> shared_from_this();  // ...
};

class Widget : public std::enable_shared_from_this<Widget> { ...
};
```

Why might a Widget want to be its own ticket?

This `run()` function passes `listener` off to `Acceptor::async_accept()`, and then immediately returns. The lambda must participate in shared ownership of the `Listener` object, to keep it from being prematurely destroyed.

```
class Listener;

void on_accept(std::shared_ptr<Listener>, std::error_code);

void run(std::shared_ptr<Listener> listener) {
    listener->acceptor().async_accept(
        listener->socket(),
        [listener](std::error_code ec) {
            on_accept(listener, ec);
        }
    );
    run(std::make_shared<Listener>(...));
}
```

Why might a Widget want to be its own ticket?

Rewrite the code in OO style. Now `Listener::run()` receives only a raw pointer — **this**. To participate in the Listener's *lifetime*, it needs to regain access to the control block. So we use `shared_from_this`.

```
class Listener : public std::enable_shared_from_this<Listener> {  
    void on_accept(std::error_code);  
    void run() {  
        acceptor_.async_accept(socket_,  
            [self = shared_from_this()](std::error_code ec) {  
                self->on_accept(ec);  
            }  
        );  
    }  
};  
std::make_shared<Listener>(...)->run();
```

In conclusion

- Use `std::unique_ptr` for unique ownership
- Use `std::shared_ptr` for reference-counting
- Don't touch raw pointers with your hands
- Pass pointers by value, not by reference
- `std::weak_ptr` is a ticket for a `shared_ptr`
- Use `enable_shared_from_this` when an object is its own ticket

Questions?